

COURSERA

Welcome to the Data Science Coding Challenge!

Test your skills in a real-world coding challenge. Coding Challenges provide CS & DS Coding Competitions with Prizes and achievement badges!

CS & DS learners want to be challenged as a way to evaluate if they're job ready. So, why not create fun challenges and give winners something truly valuable such as complimentary access to select Data Science courses, or the ability to receive an achievement badge on their Coursera Skills Profile - highlighting their performance to recruiters.

Introduction

In this challenge, you'll get the opportunity to tackle one of the most industry-relevant machine learning problems with a unique dataset that will put your modeling skills to the test. Financial loan services are leveraged by companies across many industries, from big banks to financial institutions to government loans. One of the primary objectives of companies with financial loan services is to decrease payment defaults and ensure that individuals are paying back their loans as expected. In order to do this efficiently and systematically, many companies employ machine learning to predict which individuals are at the highest risk of defaulting on their loans, so that proper interventions can be effectively deployed to the right audience.

In this challenge, we will be tackling the loan default prediction problem on a very unique and interesting group of individuals who have taken financial loans.

Imagine that you are a new data scientist at a major financial institution and you are tasked with building a model that can predict which individuals will default on their loan payments. We have provided a dataset that is a sample of individuals who received loans in 2021.

This financial institution has a vested interest in understanding the likelihood of each individual to default on their loan payments so that resources can be allocated appropriately to support these borrowers. In this challenge, you will use your machine learning toolkit to do just that!

Understanding the Datasets

Train vs. Test

In this competition, you'll gain access to two datasets that are samples of past borrowers of a financial institution that contain information about the individual and the specific loan. One dataset is titled `train.csv` and the other is titled `test.csv`.

`train.csv` contains 70% of the overall sample (255,347 borrowers to be exact) and importantly, will reveal whether or not the borrower has defaulted on their loan payments (the “ground truth”).

The `test.csv` dataset contains the exact same information about the remaining segment of the overall sample (109,435 borrowers to be exact), but does not disclose the “ground truth” for each borrower. It’s your job to predict this outcome!

Using the patterns you find in the `train.csv` data, predict whether the borrowers in `test.csv` will default on their loan payments, or not.

Dataset descriptions

Both `train.csv` and `test.csv` contain one row for each unique Loan. For each Loan, a single observation (`LoanID`) is included during which the loan was active.

In addition to this identifier column, the `train.csv` dataset also contains the target label for the task, a binary column `Default` which indicates if a borrower has defaulted on payments.

Besides that column, both datasets have an identical set of features that can be used to train your model to make predictions. Below you can see descriptions of each feature. Familiarize yourself with them so that you can harness them most effectively for this machine learning task!

```
import pandas as pd
data_descriptions = pd.read_csv('data_descriptions.csv')
pd.set_option('display.max_colwidth', None)
data_descriptions
```

	Column_name	Column_type	Data_type	\
0	LoanID	Identifier	string	
1	Age	Feature	integer	
2	Income	Feature	integer	
3	LoanAmount	Feature	integer	
4	CreditScore	Feature	integer	
5	MonthsEmployed	Feature	integer	
6	NumCreditLines	Feature	integer	
7	InterestRate	Feature	float	
8	LoanTerm	Feature	integer	
9	DTIRatio	Feature	float	
10	Education	Feature	string	
11	EmploymentType	Feature	string	
12	MaritalStatus	Feature	string	
13	HasMortgage	Feature	string	
14	HasDependents	Feature	string	
15	LoanPurpose	Feature	string	
16	HasCoSigner	Feature	string	
17	Default	Target	integer	

Description

0

A

unique identifier for each loan.
1
The age of the borrower.
2
annual income of the borrower. The
3
amount of money being borrowed. The
4
The credit score of the borrower,
indicating their creditworthiness.
5
The number of months
the borrower has been employed.
6
The number of
credit lines the borrower has open.
7
The interest rate for the loan.
8
The
term length of the loan in months.
9
The Debt-to-Income ratio, indicating the borrower's
debt compared to their income.
10
The highest level of education attained by the borrower (PhD,
Master's, Bachelor's, High School).
11
The type of employment status of the borrower (Full-time, Part-
time, Self-employed, Unemployed).
12
The marital status of the
borrower (Single, Married, Divorced).
13
Whether the
borrower has a mortgage (Yes or No).
14
Whether the
borrower has dependents (Yes or No).
15
The purpose of the loan (Home,
Auto, Education, Business, Other).
16
Whether the
loan has a co-signer (Yes or No).
17
The binary target variable indicating whether the
loan defaulted (1) or not (0).

How to Submit your Predictions to Coursera

Submission Format:

In this notebook you should follow the steps below to explore the data, train a model using the data in `train.csv`, and then score your model using the data in `test.csv`. Your final submission should be a dataframe (call it `prediction_df` with two columns and exactly 109,435 rows (plus a header row). The first column should be `LoanID` so that we know which prediction belongs to which observation. The second column should be called `predicted_probability` and should be a numeric column representing the **likelihood that the borrower will default**.

Your submission will show an error if you have extra columns (beyond `LoanID` and `predicted_probability`) or extra rows. The order of the rows does not matter.

The naming convention of the dataframe and columns are critical for our autograding, so please make sure to use the exact naming conventions of `prediction_df` with column names `LoanID` and `predicted_probability`!

To determine your final score, we will compare your `predicted_probability` predictions to the source of truth labels for the observations in `test.csv` and calculate the [ROC AUC](#). We choose this metric because we not only want to be able to predict which loans will default, but also want a well-calibrated likelihood score that can be used to target interventions and support most accurately.

Import Python Modules

First, import the primary modules that will be used in this project. Remember as this is an open-ended project please feel free to make use of any of your favorite libraries that you feel may be useful for this challenge. For example some of the following popular packages may be useful:

- pandas
- numpy
- Scipy
- Scikit-learn
- keras
- matplotlib
- seaborn
- etc, etc

```
# Import required packages

# Data packages
import pandas as pd
import numpy as np

# Machine Learning / Classification packages
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.dummy import DummyClassifier

# Visualization Packages
from matplotlib import pyplot as plt
import seaborn as sns
%matplotlib inline
```

Load the Data

Let's start by loading the dataset `train.csv` into a dataframe `train_df`, and `test.csv` into a dataframe `test_df` and display the shape of the dataframes.

```
train_df=pd.read_csv("train.csv")
print("train_df Shape:", train_df.shape)
train_df.head()
```

train_df Shape: (255347, 18)

	LoanID	Age	Income	LoanAmount	CreditScore	MonthsEmployed	\
0	I38PQUQS96	56	85994	50587	520	80	
1	HPSK72WA7R	69	50432	124440	458	15	
2	C10Z6DPJ8Y	46	84208	129188	451	26	
3	V2KKSFM3UN	32	31713	44799	743	0	
4	EY08JDHTZP	60	20437	9139	633	8	

	NumCreditLines	InterestRate	LoanTerm	DTIRatio	Education	\
0	4	15.23	36	0.44	Bachelor's	
1	1	4.81	60	0.68	Master's	
2	3	21.17	24	0.31	Master's	
3	3	7.07	24	0.23	High School	
4	4	6.51	48	0.73	Bachelor's	

	EmploymentType	MaritalStatus	HasMortgage	HasDependents	
LoanPurpose					
0	Full-time	Divorced	Yes	Yes	Other
1	Full-time	Married	No	No	Other
2	Unemployed	Divorced	Yes	Yes	Auto
3	Full-time	Married	No	No	Business
4	Unemployed	Divorced	No	Yes	Auto

	HasCoSigner	Default
0	Yes	0
1	Yes	0
2	No	1
3	No	0
4	No	0

```
test_df=pd.read_csv("test.csv")
print("test_df Shape:", test_df.shape)
test_df.head()
```

test_df Shape: (109435, 17)

	LoanID	Age	Income	LoanAmount	CreditScore	MonthsEmployed	\
0	7RYZGMKJIR	32	131645	43797	802	23	
1	JDL5RH07AM	61	134312	18402	369	87	
2	STAL716Y79	55	115809	151774	563	3	
3	S00KKJ3IQB	58	94970	55789	337	24	
4	T99CWTYDCP	63	71727	189798	451	52	

	NumCreditLines	InterestRate	LoanTerm	DTIRatio	Education	\
0	2	6.10	24	0.13	High School	
1	2	12.99	60	0.59	High School	
2	3	5.51	48	0.82	Bachelor's	
3	1	23.93	36	0.77	Bachelor's	
4	3	22.05	48	0.44	PhD	

	EmploymentType	MaritalStatus	HasMortgage	HasDependents	
0	Full-time	Divorced	Yes	No	Other
1	Self-employed	Single	No	No	Business
2	Full-time	Single	Yes	Yes	Other
3	Unemployed	Divorced	No	No	Business
4	Unemployed	Single	Yes	No	Auto

	HasCoSigner
0	No
1	Yes
2	Yes
3	No
4	No

Explore, Clean, Validate, and Visualize the Data (optional)

Feel free to explore, clean, validate, and visualize the data however you see fit for this competition to help determine or optimize your predictive model. Please note - the final autograding will only be on the accuracy of the `prediction_df` predictions.

Exploratory Data Analysis (EDA):

Before building the model, it's essential to understand data. We will perform exploratory data analysis to gain insights into the dataset. This includes checking for missing values, data types, and basic statistics.

```
train_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 255347 entries, 0 to 255346
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LoanID                255347 non-null  object
1   Age                   255347 non-null  int64
2   Income                255347 non-null  int64
3   LoanAmount            255347 non-null  int64
4   CreditScore           255347 non-null  int64
5   MonthsEmployed        255347 non-null  int64
6   NumCreditLines        255347 non-null  int64
7   InterestRate          255347 non-null  float64
8   LoanTerm              255347 non-null  int64
9   DTIRatio              255347 non-null  float64
10  Education              255347 non-null  object
11  EmploymentType        255347 non-null  object
12  MaritalStatus         255347 non-null  object
13  HasMortgage           255347 non-null  object
14  HasDependents         255347 non-null  object
15  LoanPurpose           255347 non-null  object
16  HasCoSigner           255347 non-null  object
17  Default               255347 non-null  int64
dtypes: float64(2), int64(8), object(8)
memory usage: 35.1+ MB

```

```
train_df.isnull().sum()
```

```

LoanID      0
Age         0
Income      0
LoanAmount  0
CreditScore 0
MonthsEmployed 0
NumCreditLines 0
InterestRate 0
LoanTerm    0
DTIRatio    0
Education   0
EmploymentType 0
MaritalStatus 0
HasMortgage 0
HasDependents 0
LoanPurpose 0
HasCoSigner 0
Default     0
dtype: int64

```

```
test_df.isnull().sum()
```

LoanID	0
Age	0
Income	0
LoanAmount	0
CreditScore	0
MonthsEmployed	0
NumCreditLines	0
InterestRate	0
LoanTerm	0
DTIRatio	0
Education	0
EmploymentType	0
MaritalStatus	0
HasMortgage	0
HasDependents	0
LoanPurpose	0
HasCoSigner	0
dtype: int64	

Both `train.csv` and `test.csv` have no missing values and datatypes are identical

Data Preprocessing:

Prepare the data for training by handling missing values, encoding categorical variables, and splitting the training data into features (X) and the target (y).

```
# Convert categorical columns to numeric using Label Encoding
categorical_columns=['Education', 'EmploymentType', 'MaritalStatus',
                    'HasMortgage', 'HasDependents', 'LoanPurpose', 'HasCoSigner']
label_encoders={}

for column in categorical_columns:
    label_encoder=LabelEncoder()
    train_df[column]=label_encoder.fit_transform(train_df[column])
    test_df[column]=label_encoder.fit_transform(test_df[column])
    label_encoders[column]=label_encoder

train_df.head()
```

	LoanID	Age	Income	LoanAmount	CreditScore	MonthsEmployed	\
0	I38PQUQS96	56	85994	50587	520	80	
1	HPSK72WA7R	69	50432	124440	458	15	
2	C10Z6DPJ8Y	46	84208	129188	451	26	
3	V2KKSFM3UN	32	31713	44799	743	0	
4	EY08JDHTZP	60	20437	9139	633	8	

	NumCreditLines	InterestRate	LoanTerm	DTIRatio	Education	\
0	4	15.23	36	0.44	0	

1	1	4.81	60	0.68	2
2	3	21.17	24	0.31	2
3	3	7.07	24	0.23	1
4	4	6.51	48	0.73	0

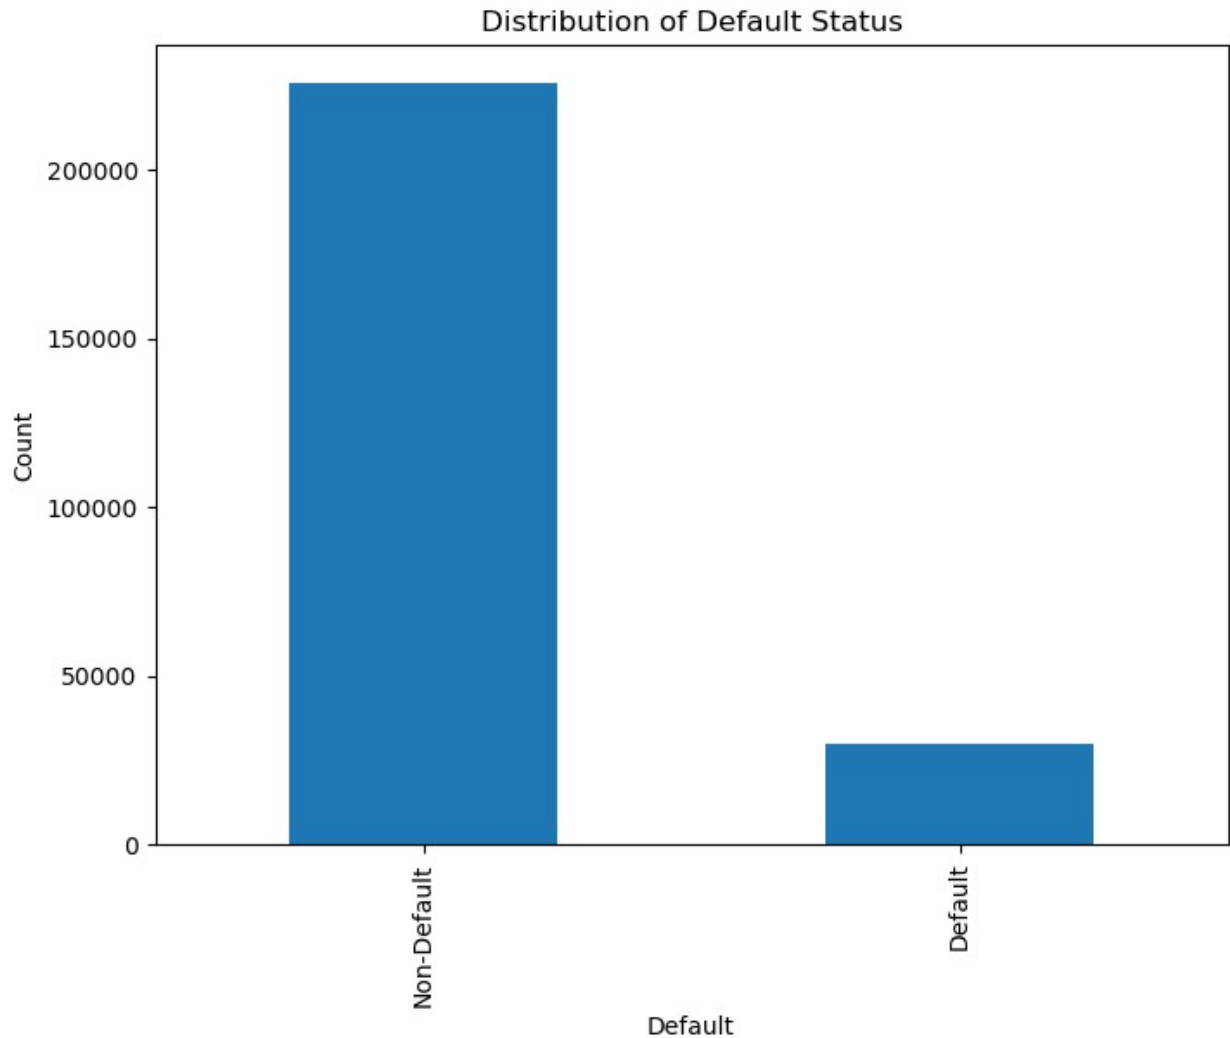
LoanPurpose \	EmploymentType	MaritalStatus	HasMortgage	HasDependents
0	0	0	1	1
4				
1	0	1	0	0
4				
2	3	0	1	1
0				
3	0	1	0	0
1				
4	3	0	0	1
0				

HasCoSigner	Default
0	1
1	1
2	0
3	0
4	0

Histogram of Default Status

This histogram will show the distribution of default and non-default loans in the training data.

```
# Plot a histogram of Default status
plt.figure(figsize=(8, 6))
train_df['Default'].value_counts().plot(kind='bar')
plt.title('Distribution of Default Status')
plt.xlabel('Default')
plt.ylabel('Count')
plt.xticks([0, 1], ['Non-Default', 'Default'])
plt.show();
```



Correlation Matrix:

A correlation matrix will help us understand the relationships between different features

Calculation of the correlation matrix

```
correlation_matrix=train_df.corr()
correlation_matrix
```

	Age	Income	LoanAmount	CreditScore	
MonthsEmployed	\				
Age	1.000000	-0.001244	-0.002213	-0.000548	-
0.000341					
Income	-0.001244	1.000000	-0.000865	-0.001430	
0.002675					
LoanAmount	-0.002213	-0.000865	1.000000	0.001261	
0.002817					
CreditScore	-0.000548	-0.001430	0.001261	1.000000	
0.000613					

MonthsEmployed	-0.000341	0.002675	0.002817	0.000613	
1.000000					
NumCreditLines	-0.000890	-0.002016	0.000794	0.000016	
0.001267					
InterestRate	-0.001127	-0.002303	-0.002291	0.000436	
0.000096					
LoanTerm	0.000263	-0.000998	0.002538	0.001130	-
0.001166					
DTIRatio	-0.004689	0.000205	0.001122	-0.001039	
0.001765					
Education	-0.000882	-0.000965	0.002551	0.000214	-
0.001304					
EmploymentType	0.000787	-0.005146	0.003060	0.003503	
0.000564					
MaritalStatus	-0.002187	0.000637	-0.000771	-0.003218	-
0.000095					
HasMortgage	0.000035	-0.000945	-0.000801	0.001728	
0.000210					
HasDependents	0.000710	-0.001570	0.000139	-0.003018	
0.001450					
LoanPurpose	0.002264	-0.002092	0.000057	0.000596	-
0.002579					
HasCoSigner	-0.002918	-0.003524	-0.001848	-0.002755	
0.001045					
Default	-0.167783	-0.099119	0.086659	-0.034166	-
0.097374					

	NumCreditLines	InterestRate	LoanTerm	DTIRatio	
Education \					
Age	-0.000890	-0.001127	0.000263	-0.004689	-
0.000882					
Income	-0.002016	-0.002303	-0.000998	0.000205	-
0.000965					
LoanAmount	0.000794	-0.002291	0.002538	0.001122	
0.002551					
CreditScore	0.000016	0.000436	0.001130	-0.001039	
0.000214					
MonthsEmployed	0.001267	0.000096	-0.001166	0.001765	-
0.001304					
NumCreditLines	1.000000	-0.000297	-0.000226	-0.000586	
0.002691					
InterestRate	-0.000297	1.000000	0.000892	0.000575	
0.002879					
LoanTerm	-0.000226	0.000892	1.000000	0.002273	-
0.002999					
DTIRatio	-0.000586	0.000575	0.002273	1.000000	
0.001789					
Education	0.002691	0.002879	-0.002999	0.001789	
1.000000					

EmploymentType 0.000236	0.000219	0.000525	0.000779	-0.000578	
MaritalStatus 0.004717	-0.000664	-0.005079	-0.001042	0.004492	-
HasMortgage 0.001167	-0.001744	-0.000424	0.001775	0.000231	
HasDependents 0.001048	-0.001895	-0.000243	0.002417	0.001492	
LoanPurpose 0.003271	0.000340	0.001472	0.002856	-0.003819	-
HasCoSigner 0.001707	0.002105	-0.003991	-0.001166	0.000373	
Default 0.022835	0.028330	0.131273	0.000545	0.019236	-


	EmploymentType	MaritalStatus	HasMortgage	
HasDependents \				
Age 0.000710	0.000787	-0.002187	0.000035	
Income 0.001570	-0.005146	0.000637	-0.000945	-
LoanAmount 0.000139	0.003060	-0.000771	-0.000801	
CreditScore 0.003018	0.003503	-0.003218	0.001728	-
MonthsEmployed 0.001450	0.000564	-0.000095	0.000210	
NumCreditLines 0.001895	0.000219	-0.000664	-0.001744	-
InterestRate 0.000243	0.000525	-0.005079	-0.000424	-
LoanTerm 0.002417	0.000779	-0.001042	0.001775	
DTIRatio 0.001492	-0.000578	0.004492	0.000231	
Education 0.001048	0.000236	-0.004717	0.001167	
EmploymentType 0.002480	1.000000	0.002768	0.001193	
MaritalStatus 0.000437	0.002768	1.000000	-0.000408	-
HasMortgage 0.000067	0.001193	-0.000408	1.000000	
HasDependents 1.000000	0.002480	-0.000437	0.000067	
LoanPurpose 0.003759	0.000734	0.001434	-0.002157	-
HasCoSigner 0.001602	-0.000033	-0.000888	-0.003529	

Default	0.041010	-0.007902	-0.022856	-
0.034678				

	LoanPurpose	HasCoSigner	Default
Age	0.002264	-0.002918	-0.167783
Income	-0.002092	-0.003524	-0.099119
LoanAmount	0.000057	-0.001848	0.086659
CreditScore	0.000596	-0.002755	-0.034166
MonthsEmployed	-0.002579	0.001045	-0.097374
NumCreditLines	0.000340	0.002105	0.028330
InterestRate	0.001472	-0.003991	0.131273
LoanTerm	0.002856	-0.001166	0.000545
DTIRatio	-0.003819	0.000373	0.019236
Education	-0.003271	0.001707	-0.022835
EmploymentType	0.000734	-0.000033	0.041010
MaritalStatus	0.001434	-0.000888	-0.007902
HasMortgage	-0.002157	-0.003529	-0.022856
HasDependents	-0.003759	0.001602	-0.034678
LoanPurpose	1.000000	-0.001935	-0.010096
HasCoSigner	-0.001935	1.000000	-0.039109
Default	-0.010096	-0.039109	1.000000

```
# Create a heatmap of the correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f",
            cmap='coolwarm', linewidths=0.5, square=True)

# Adjust font size and rotation for the labels
plt.xticks(fontsize=10, rotation=45)
plt.yticks(fontsize=10)
plt.title('Correlation Matrix', fontsize=12)
plt.show()
```



Train a machine learning model

```
# Define features and target variable
X=train_df.drop(['LoanID', 'Default'], axis=1)
y=train_df['Default']

# Split the data into a training set and a validation set
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
random_state=42)

# Initialize and train a Random Forest Classifier
model=RandomForestClassifier()
model.fit(X, y)

RandomForestClassifier()
```

```

model.score(X_train, y_train)

0.9999804187451353

# Predict Probabilities
predicted_probabilities=model.predict_proba(X_val)[:, 1]

```

Calculate ROC AUC

To determine your final score, we will compare your predicted_probability predictions to the source of truth labels for the observations

```

roc_auc=roc_auc_score(y_val, predicted_probabilities)

print(f'Validation ROC AUC Score: {roc_auc}')

Validation ROC AUC Score: 1.0

# Now, we can use the trained model to make predictions on the test data
X_test=test_df.drop(columns=['LoanID'])

# Predict probabilities for the test data
predicted_probabilities=model.predict_proba(X_test)[:, 1]

```

Create the submission dataframe with LoanID and predicted_probability and save it.

```

# Predicted_probabilities probabilities should be 0 and 1

threshold = 0.5
binary_predictions=(predicted_probabilities>=threshold).astype(int)

prediction_df=pd.DataFrame({
    'LoanID': test_df[['LoanID']].values[:, 0],
    'predicted_probability': binary_predictions
})

prediction_df.head(10)

```

	LoanID	predicted_probability
0	7RYZGMKJIR	0
1	JDL5RH07AM	0
2	STAL716Y79	0
3	S00KKJ3IQB	0
4	T99CWTYDCP	0
5	0SNHFWV4UP	0
6	S6ITP6LGYS	0
7	A6I7U12IRJ	0

```

8  8W6KY50JU4      0
9  THFQ080LMU      0

prediction_df.shape

(109435, 2)

prediction_df.predicted_probability.value_counts()

0      108299
1       1136
Name: predicted_probability, dtype: int64

```

Final Tests - **IMPORTANT** - the cells below must be run prior to submission

Below are some tests to ensure your submission is in the correct format for autograding. The autograding process accepts a csv `prediction_submission.csv` which we will generate from our `prediction_df` below. Please run the tests below and ensure no assertion errors are thrown.

```

# FINAL TEST CELLS - please make sure all of your code is above these
test cells

# Writing to csv for autograding purposes
prediction_df.to_csv("prediction_submission.csv", index=False)
submission = pd.read_csv("prediction_submission.csv")

assert isinstance(submission, pd.DataFrame), 'You should have a
dataframe named prediction_df.'

# FINAL TEST CELLS - please make sure all of your code is above these
test cells

assert submission.columns[0] == 'LoanID', 'The first column name
should be CustomerID.'
assert submission.columns[1] == 'predicted_probability', 'The second
column name should be predicted_probability.'

# FINAL TEST CELLS - please make sure all of your code is above these
test cells

assert submission.shape[0] == 109435, 'The dataframe prediction_df
should have 109435 rows.'

# FINAL TEST CELLS - please make sure all of your code is above these
test cells

assert submission.shape[1] == 2, 'The dataframe prediction_df should
have 2 columns.'

```



```
roc_auc=roc_auc_score(y_val, predicted_probabilities)
```

```
print(f'Validation ROC AUC Score: {roc_auc}')
```

```
Validation ROC AUC Score: 1.0
```